



Parvatibai Genba Moze College of Engineering

Department of Computer Engineering

Lab Manual

Laboratory Practice V

BE Computer (2019 Pattern) Subject Code: 410255

Academic Year: 2023-2024

Savitribai Phule Pune University Fourth Year of Computer Engineering (2019 Course) 410255: Laboratory Practice V

Teaching Scheme: PR: 02 Credi Examination Scheme: Hours/Week t01 Examination Scheme: Term Work: 50 Marks Practical: 50 Marks

Guidelines for Instructor's Manual

Laboratory Practice V is for practical hands on for core courses High Performance Computing and Data Learning. The instructor's manual is to be developed as a hands-on resource and as ready reference. The instructor's manual need to include prologue (about university/program/ institute/ department/foreword/ preface etc.), University syllabus, conduction and Assessment guidelines, topics under consideration-concept, objectives, outcomes, set of typical applications/assignments/ guidelines, references among others.

Guidelines for Student's Laboratory Journal

The laboratory assignments are to be submitted by student in the form of journal. Journal may consist of prologue, Certificate, table of contents, and handwritten write-up of each assignment (Title, Objectives, Problem Statement, Outcomes, software and Hardware requirements, Date of Completion, Assessment grade/marks and assessor's sign, Theory- Concept in brief, Algorithm/Database design, test cases, conclusion/analysis). Program codes with sample output of all performed assignments are to be submitted as softcopy.

Guidelines for Laboratory /TW Assessment

Continuous assessment of laboratory work is to be done based on overall performance and lab assignments performance of student. Each lab assignment assessment will assign grade/marks based on parameters with appropriate weightage. Suggested parameters for overall assessment as well as each lab assignment assessment include- timely completion, performance, innovation, efficient codes, punctuality and neatness reserving weightage for successful mini-project completion and related documentation

Guidelines for Laboratory Conduction

- List of recommended programming assignments and sample mini-projects is provided for reference.
- Referring these, Course Teacher or Lab Instructor may frame the assignments/mini-project by understanding the prerequisites, technological aspects, utility and recent trends related to the respective courses.
- Preferably there should be multiple sets of assignments/mini-project and distribute among batches of students.
- Real world problems/application-based assignments/mini-projects create interest among learners serving as foundation for future research or startup of business projects.
- Mini-project can be completed in group of 2 to 3 students.
- Software Engineering approach with proper documentation is to be strictly followed.
- Use of open-source software is to be encouraged.
- Instructor may also set one assignment or mini-project that is suitable to respective course beyond the scope of syllabus.

Operating System recommended: - 64-bit Open-source Linux or its derivative

Programming Languages: - Object Oriented Languages C++/JAVA/PYTHON/R.

Programming tools recommended: Front End: Java/Perl/PHP/Python/Ruby/.net,

Backend: MongoDB/MYSQL/Oracle, Database Connectivity: ODBC/JDBC

LABORATORY EXPERIMENTS

| | 410250: High Performance Computing |
|----|--|
| | |
| 5 | |
| | GROUP 1 |
| 1. | Design and implement Parallel Breadth First Search and Depth First Search based on existing algorithms using OpenMP. Use a Tree or an undirected graph for BFS and DFS |
| 2. | Write a program to implement Parallel Bubble Sort and Merge sort using OpenMP. Use existing algorithms and measure the performance of sequential and parallel algorithms |
| 3 | Implement Min, Max, Sum and Average operations using Parallel Reduction. |
| 4 | Write a CUDA Program for: 1. Addition of two large vectors 2. Matrix Multiplication using CUDA C |
| 5 | Implement HPC application for AI/ML domain. |
| | GROUP 2 |
| 6 | Mini Project: Evaluate performance enhancement of parallel Quicksort Algorithm using MPI. |
| 7 | Mini Project: Implement Huffman Encoding on GPU. |
| 8 | Mini Project: Implement Parallelization of Database Query optimization. |
| 9 | Mini Project: Implement Non-Serial Polyadic Dynamic Programming with GPU Parallelization |

| | 410251: Deep Learning |
|---|--|
| | Group 1 |
| 1 | Linear regression by using Deep Neural network: Implement Boston housing price prediction problem by Linear regression using Deep Neural network. Use Boston House price prediction dataset. |
| 2 | Classification using Deep neural network (Any One from the following) 1. Multiclass classification using Deep Neural Networks: Example: Use the OCR letter recognition datasethttps://archive.ics.uci.edu/ml/datasets/letter+recognition 2. Binary classification using Deep Neural Networks Example: Classify movie reviews into positive" reviews and "negative" reviews, just based on the text content of the reviews. Use IMDB dataset |
| 3 | Convolutional neural network (CNN) (Any One from the following) Use any dataset of plant disease and design a plant disease detection system using CNN. Use MNIST Fashion Dataset and create a classifier to classify fashion clothing into categories. |
| 4 | Recurrent neural network (RNN) Use the Google stock prices dataset and design a time series analysis and prediction system using RNN |
| | Group 2 |
| 5 | Mini Project: Human Face Recognition |
| 6 | Mini Project: Gender and Age Detection: predict if a person is a male or female and also their age |
| 7 | Mini Project: Colorizing Old B&W Images: color old black and white images to colorful images |

Assignment No 1

<u>Title</u>: Design and implement Parallel Breadth First Search and Depth First Search based on existing algorithms using OpenMP.

Objectives: Use a Tree or an undirected graph for BFS and DFS.

Theory:

Breadth First Search

To design and implement parallel breadth first search using OpenMP, you can use the existing breadth first search algorithm and parallelize it using OpenMP's palatalization constructs.

Parallel Breadth First Search

- 1. To design and implement parallel breadth first search, you will need to divide the graph into smaller sub-graphs and assign each sub-graph to a different processor or thread.
- 2. Each processor or thread will then perform a breadth first search on its assigned sub-graph concurrently with the other processors or threads.
- 3. Two methods: Vertex by Vertex OR Level by Level

```
Procedure
              Parallel-Breadth-First-Search-Vertex(ALM, EM, U)
begin
     mark every vertex "unvisited"
     v ← start vertex
     mark v "visited"
     instruct processor(i) where 1 \le i \le k
          for j = 1 to k do
               if (k * (j-1) + i) \leq EM(v)
               then delete v from U(ALM(v, k * (j-1) + i))
          endfor
     end-instruction
     initialize queue with v
     while queue is not empty do
          begin
          v ← first vertex from the queue
          for each w∈ U(v) do
               begin
               mark w "visited"
               instruct processor (i) where 1 \le i \le k
                   for j = 1 to k do
                              if (k * (j-1) + i) \le EM(w)
                               then delete w from U(ALM(w, k * (j - 1) + i))
                     endfor
                end-instruction
                add w to queue
                end
            endfor
       endwhile
  end
```

Parallel bfs steps

- 1. In this implementation, the parallel bfs function takes in a graph represented as an adjacency list, where each element in the list is a vector of neighbouring vertices, and a starting vertex.
- 2. The bfs function uses a queue to keep track of the vertices to visit, and a Boolean visited array to keep track of which vertices have been visited. The #pragma omp parallel directive creates a parallel region and the #pragma omp single directive creates a single execution context within that region.
- 3. Inside the while loop, the #pragma omp task directive creates a new task for each unvisited neighbour of the current vertex.
- 4. This allows each task to be executed in parallel with other tasks. The first private clause is used to ensure that each task has its own copy of the vertex variable.
- 5. This is just one possible implementation, and there are many ways to improve it depending on the specific requirements of your application. For example, you can use omp atomic or omp critical to protect the shared resource queue.
- 6. This is just one possible implementation, and there are many ways to improve it depending on the specific requirements of your application. For example, you can use omp atomic or omp critical to protect the shared resource queue.

Parallel Depth-First Search

- Different subtrees can be searched concurrently.
- Subtrees can be very different in size.
- Estimate the size of a subtree rooted at a node.
- Dynamic load balancing is required.

Parallel Depth-First Search: Dynamic Load Balancing

- When a processor runs out of work, it gets more work from another processor.
- This is done using work requests and responses in message passing machines and locking and extracting work in shared address space machines.
- On reaching final state at a processor, all processors terminate.

- Unexplored states can be conveniently stored as local stacks at processors.
- The entire space is assigned to one processor to begin with.

Parameters in Parallel DFS: Work Splitting

- Work is split by splitting the stack into two.
- Ideally, we do not want either of the split pieces to be small.
- Select nodes near the bottom of the stack (node splitting), or
- Select some nodes from each level (stack splitting).
- The second strategy generally yields a more even split of the space.

Load-Balancing Schemes

- Asynchronous round robin: Each processor maintains a counter and makes requests in a round- robin fashion.
- Global round robin: The system maintains a global counter and requests are made in a round- robin fashion, globally.
- Random polling: Request a randomly selected processor for work.

Analyzing DFS

- We can't compute, analytically, the serial work W or parallel time. Instead, we quantify total overhead T_O in terms of W to compute scalability.
- For dynamic load balancing, idling time is subsumed by communication.
- We must quantify the total number of requests in the system.

Load-Balancing Schemes

- Asynchronous round robin has poor performance because it makes a large number of work requests.
- Global round robin has poor performance because of contention at counter, although it
 makes the least number of requests.
- Random polling strikes a desirable compromise.

Termination Detection

- Processor P0 has all the work and a weight of one is associated with it. When its work
 is partitioned and sent to another processor, processor P0 retains half of the weight and
 gives half of it to the processor receiving the work.
- If Pi is the recipient processor and wi is the weight at processor Pi, then after the first work transfer, both w0 and wi are 0.5.
- Each time the work at a processor is partitioned, the weight is halved. When a processor completes its computation, it returns its weight to the processor from which it received work.
- Termination is signalled when the weight w0 at processor P0 becomes one and processor P0 has finished its work

```
Program
#include <iostream>
#include <vector>
#include <stack>
#include <omp.h>
using namespace std;
void dfs(vector<vector<int>>& graph, int start,
vector<bool>& visited) {
  stack<int> s:
  s.push(start);
  visited[start] = true;
#pragma omp parallel
     #pragma omp single
       while (!s.empty()) {
          int vertex = s.top();
          s.pop();
```

Parellel dfs Steps:

- 1. In this implementation, the parallel dfs function takes in a graph represented as an adjacency list, where each element in the list is a vector of neighboring vertices, and a starting vertex.
- 2. The dfs function uses a stack to keep track of the vertices to visit, and a boolean visited array to keep track of which vertices have been visited.
- 3. The #pragma omp parallel directive creates a parallel region and the #pragma omp single directive creates a single execution context within that region.
- 4. Inside the while loop, the #pragma omp task directive creates a new task for each unvisited neighbor of the current vertex.

- 5. This allows each task to be executed in parallel with other tasks. The first private clause is used to ensure that each task has its own copy of the vertex variable.
- 6. This implementation is suitable for both tree and undirected graph, since both are represented as an adjacency list and the algorithm is using a stack to traverse the graph.
- 7. This is just one possible implementation, and there are many ways to improve it depending on the specific requirements of your application. For example, you can use omp atomic or omp critical to protect the shared resource stack.
- 8. The dfs function uses a stack to keep track of the vertices to visit, and a boolean visited array to keep track of which vertices have been visited.
- 9. The #pragma omp parallel directive creates a parallel region and the #pragma omp single directive creates a single execution context within that region.

Conclusion: Here we implement the parallel dfs and bfs algorithm.

Questions:

Assignment No 2

<u>Title</u>: Write a program to implement Parallel Bubble Sort and Merge sort using OpenMP. Use existing algorithms and measure the performance of sequential and parallel algorithms.

Objectives:

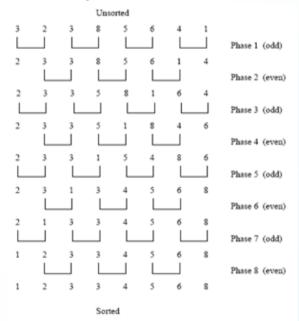
- 1. Implement parallel bubble and merge sort.
- 2. measuring the performance using sequential and parallel algorithm.

Theory:

BUBBLE SORT

- 1. The complexity of bubble sort is $\Theta(n^2)$.
- 2. Bubble sort is difficult to parallelize since the algorithm has no concurrency.
- 3. A simple variant, though, uncovers the concurrency.

Odd-Even Transposition



Sorting n = 8 elements, using the odd-even transposition sort algorithm. During each phase, n = 8 elements are compared

Odd-Even Transposition

```
procedure ODD-EVEN(n)
1.
2.
         begin
3.
              for i := 1 to n do
4.
              begin
                   if i is odd then
5.
                        for j := 0 to n/2 - 1 do
6.
7.
                            compare-exchange(a_{2j+1}, a_{2j+2});
8.
                   if i is even then
9.
                        for j := 1 to n/2 - 1 do
                            compare-exchange (a_{2i}, a_{2i+1});
10.
11.
              end for
         end ODD-EVEN
12.
```

Sequential odd-even transposition sort algorithm.

Odd-Even Transposition

- 1. After *n* phases of odd-even exchanges, the sequence is sorted.
- 2. Each phase of the algorithm (either odd or even) requires $\Theta(n)$ comparisons.
- 3. Serial complexity is $\Theta(n^2)$.

Parallel Odd-Even Transposition

- 1. Consider the one item per processor case.
- 2. There are n iterations, in each iteration, each processor does one compare-exchange.
- 3. The parallel run time of this formulation is $\Theta(n)$.
- 4. This is cost optimal with respect to the base serial algorithm but not the optimal one.

Parallel Odd-Even Transposition

```
procedure ODD-EVEN_PAR(n)
        begin
            id := process's label
4.
            for i := 1 to n do
            begin
                 if i is odd then
7.
                     if id is odd then
                         compare-exchange_min(id + 1);
                     else
                         compare-exchange_max(id - 1);
11.
                 if i is even then
                     if id is even then
12.
13.
                         compare-exchange_min(id + 1);
14.
                     else
15.
                         compare-exchange_max(id - 1);
16.
            end for
        end ODD-EVEN_PAR
17.
```

Parallel formulation of odd-even transposition.

Parallel Odd-Even Transposition

- 1. Consider a block of n/p elements per processor.
- 2. The first step is a local sort.
- 3. In each subsequent step, the compare exchange operation is replaced by the compare split operation.
- 4. The parallel run time of the formulation is

$$T_P = \overbrace{\Theta\left(\frac{n}{p}\log\frac{n}{p}\right)}^{\text{local sort}} + \overbrace{\Theta(n)}^{\text{comparisons}} + \overbrace{\Theta(n)}^{\text{communication}}.$$

Bubble Sort Odd Even Transposition

```
#include <iostream>
#include <vector>
#include <omp.h>
using namespace std;
void bubble sort odd even(vector<int>& arr) {
  bool isSorted = false:
  while (!isSorted) {
     isSorted = true:
     #pragma omp parallel for
     for (int i = 0; i < arr.size() - 1; i += 2) {
       if (arr[i] > arr[i+1]) {
          swap(arr[i], arr[i+1]);
          isSorted = false:
     #pragma omp parallel for
     for (int i = 1; i < arr.size() - 1; i += 2) {
       if (arr[i] > arr[i+1]) {
          swap(arr[i], arr[i+1]);
          isSorted = false;
```

```
int main () {
   vector<int> arr = {5, 2, 9, 1, 7, 6, 8, 3, 4};
   double start, end;

// Measure performance of parallel bubble sort using odd- even transposition
   start = omp_get_wtime ();
   bubble_sort_odd_even(arr); end =
   omp_get_wtime ();
   cout << "Parallel bubble sort using odd-even transposition time: " << end - start
   << endl;
}</pre>
```

Performance of sequential and parallel algorithms

- 1. This program uses OpenMP to parallelize the bubble sort algorithm.
- 2. The **#pragma omp parallel for** directive tells the compiler to create a team of threads to execute the for loop within the block in parallel.
- 3. Each thread will work on a different iteration of the loop, in this case on comparing and swapping the elements of the array.
- 4. The **bubbleSort** function takes in an array, and it sorts it using the bubble sort algorithm. The outer loop iterates from 0 to n-2 and the inner loop iterates from 0 to n-i-1, where i is the index of the outer loop. The inner loop compares the current element with the next element, and if the current element is greater than the next element, they are swapped.
- 5. The **main** function creates a sample array and calls the **bubbleSort** function to sort it. The sorted array is then printed.
- 6. This is a skeleton code and it may not run as is and may need some modification to work with specific inputs and requirements.
- 7. It is worth noting that bubble sort is not an efficient sorting algorithm, specially for large inputs, and it may not scale well with more number of threads. Also parallelizing bubble sort does not have a significant improvement in performance due to the nature of the algorithm itself.
- 8. In this implementation, the bubble_sort_odd_even function takes in an array and sorts it using the odd-even transposition algorithm. The outer while loop continues until the array is sorted. Inside the loop, the #pragma omp parallel for directive creates a parallel region and divides the loop iterations among the available threads. Each

- thread performs the swap operation in parallel, improving the performance of the algorithm.
- 9. The two #pragma omp parallel for inside while loop, one for even indexes and one for odd indexes, allows each thread to sort the even and odd indexed elements simultaneously and prevent the dependency.

MERGE SORT

Parallel Merge Sort

Given a set of elements $A = \{a ... \sim, \sim, ..., an\}$, Aodd and Aeven are defined as the set of elements of A with odd and even indices, respectively.

For example, Aodd = $\{a ... a3, a" ... \}$ and Aeven = $\{ \sim, a4, a6" ... \}$ regarding a set of elements A = $\{a ... an\}$ ' Similarly, let a set of elements B = $\{b bn\}$. We can then define the merge operation as:

$$\begin{split} & Merge(A,B) = \{a.,b.,\sim,b2,a3,b3, \dots,an,bn\} \\ & For \ example, \\ & \text{if} \ A = \{1,2,3,4\} \ and \ B = \{5,6,7,8\} \\ & \text{then} \ Merge(\ (1,2,3,4\},\{5,6,7,8\}) = \{1,5,2,6,3,7,4,8\} \end{split}$$

Join(A,B) = (Merge(A,B), Odd-Even(A,B))

```
Algorithm: Odd-Even(A,B,S)
      begin
      if A and B are of length 1
            then
 Merge A and B using one Compare-and-Exchange operation
            else
            begin
             compute Sodd and Seven In Parallel do
             S_{odd} = Merge(A_{odd}, B_{odd})
             S_{even} = Merge(A_{even}, B_{even})
             S_{odd-even} = Join(S_{odd}, S_{even})
             end
       endif
       end
Example: Suppose the set of elements S =
\{2,3,6,10,15,4,5,8\} and we start with A =
\{2,6,IO,15\} and B = \{3,4,5,8\}, two sorted sets
of elements. Then
                             Merge(A_{out}, B_{out}) = \{2,3,5,10\}
                             Merge(A_{even}, B_{even}) = \{4,6,8,15\}
               The join operation:
                       Join(A,B) = \{Merge(A,B),Odd-Even(A,B)\}\
```

requires a merge operation, which results in $Merge(A,B) = \{2,4,3,6,5,8,10,15\},\$

and an odd-even operation, which obtains the final sorted list of elements,

Odd-Even $\{2,4,3,6,5,8,10,15\} = \{2,3,4,5,6,8,10,15\}$

```
#include < iostream > #include < vector > #include < omp.h >
  using namespace std;
  void merge(vector\leqint\geq& arr, int l, int m, int r) { int i, j, k;
     int n1 = m - 1 + 1; int n2 = r - m;
     vector < int > L(n1), R(n2);
     for (i = 0; i < n1; i++) \{ L[i] = arr[1+i];
     for (j = 0; j < n2; j++) { R[j] = arr[m+1+j];
     i = 0;
     j = 0;
     k = 1:
     while (i < n1 \&\& j < n2) \{ if (L[i] \le R[j]) \}
            arr[k++] = L[i++];
         } else {
            arr[k++] = R[j++];
         }
     }
void merge sort(vector<int>& arr, int l, int r) {
 if (1 \le r)
   int m = 1 + (r - 1)/2; #pragma omp task merge sort(arr, 1, m); #pragma omp task merge sort(arr, m + 1, r); merge(arr, 1, m, r);
void parallel merge sort(vector<int>& arr) {
 #pragma omp parallel
   #pragma omp single merge sort(arr, 0, arr.size() - 1);
```

```
int main() {
   vector<int> arr = {5, 2, 9, 1, 7, 6, 8, 3, 4}; double start, end;

// Measure performance of sequential merge sort start = omp_get_wtime();
   merge_sort(arr, 0, arr.size() - 1); end = omp_get_wtime();
   cout << "Sequential merge sort time: " << end - start <<
endl;

// Measure performance of parallel merge sort arr = {5, 2, 9, 1, 7, 6, 8, 3, 4};
   start = omp_get_wtime(); parallel merge_sort(arr); end = omp_get_</pre>
```

Conclusion:

Hence, we studied the parallel bubble sort and merge sort algorithm and check the performance of sequential and parallel algorithm.

Questions:

Assignment No 3

Title: Implement Min, Max, Sum and Average operations using Parallel Reduction

Objectives: Perform the Min, Max Sum and Average operation using parallel reduction technique.

Theory:

The min reduction function finds the minimum value in the input array using the #pragma omp parallel for reduction(min: min_value) directive, which creates a parallel region and divides the loop iterations among the available threads. Each thread performs the comparison operation in parallel and updates the min_value variable if a smaller value is found.

Similarly, the max_reduction function finds the maximum value in the array, sum_reduction function finds the sum of the elements of array and average_reduction function finds the average of the elements of array by dividing the sum by the size of the array.

The reduction clause is used to combine the results of multiple threads into a single value, which is then returned by the function. The min and max operators are used for the min_reduction and max_reduction functions, respectively, and the + operator is used for the sum_reduction and average_reduction functions. In the main function, it creates a vector and calls the functions min_reduction, max_reduction, sum_reduction, and average_reduction to compute the values of min, max, sum and average respectively.

```
#include <iostream> #include <vector> #include <omp.h>

using namespace std;
void min_reduction(vector<int>& arr) {
    int min_value = INT_MAX;
    #pragma omp parallel for reduction(min: min_value) for (int i = 0; i < arr.size(); i++) {
        if (arr[i] < min_value) { min_value = arr[i];
        }
    }
    cout << "Minimum value: " << min_value << endl;
}

void max_reduction(vector<int>& arr) { int max_value = INT_MIN;
    #pragma omp parallel for reduction(max: max_value)
    for (int i = 0; i < arr.size(); i++) { if (arr[i] > max_value) {
            max_value = arr[i];
        }
    }
    cout << "Maximum value: " << max_value << endl;
}</pre>
```

```
int sum = 0;
    #pragma omp parallel for reduction(+: sum) for (int i = 0; i < arr.size(); i++) {
        sum += arr[i];
    }
    cout << "Sum: " << sum << endl;
}
void average_reduction(vector<int>& arr) { int sum = 0;
    #pragma omp parallel for reduction(+: sum) for (int i = 0; i < arr.size(); i++) {
        sum += arr[i];
    }
    cout << "Average: " << (double)sum / arr.size() << endl;
}
int main() {
    vector<int> arr = {5, 2, 9, 1, 7, 6, 8, 3, 4};
    min_reduction(arr); max_reduction(arr);
    sum_reduction(arr); average_reduction(arr);
```

Conclusion:

Hence, we Implement Min, Max, Sum and Average operations using Parallel Reduction

Questions:

Assignment No 4

<u>Title</u>: Write a CUDA Program for:

- 1. Addition of two large vectors
- 2. Matrix Multiplication using CUDA C

Objectives: Perform the addition operation of two vectors and Multiplication of matrix suing CUDA C.

Theory:

In this program, the `addVectors` kernel takes in the two input vectors `A` and `B`, the output vector `C`, and the size of the vectors `n`. The kernel uses the `blockIdx.x` and `threadIdx.x` variables to calculate the index `i` of the current thread. If the index is less than `n`, the kernel performs the addition operation C[i] = A[i] + B[i]`.

In the `main` function, the program first allocates memory for the input and output vectors on the host and initializes them. Then it allocates memory for the vectors on the device and copies the data from the host to the device using `cudaMemcpy`.

```
#include <iostream>
#include <cuda runtime.h>
 _global__void addVectors(int*A, int*B, int*C, int n) {
  int i = blockIdx.x * blockDim.x + threadIdx.x;
  if (i < n) {
    C[i] = A[i] + B[i];
int main() {
  int n = 1000000;
  int* A, * B, * C;
  int size = n * sizeof(int);
  // Allocate memory on the host
  cudaMallocHost(&A, size);
  cudaMallocHost(&B, size);
  cudaMallocHost(&C, size);
  // Initialize the vectors
  for (int i = 0; i < n; i++) {
    A[i] = i;
    B[i] = i * 2;
```

```
// Allocate memory on the device
  int* dev A, * dev B, * dev C;
  cudaMalloc(&dev A, size);
  cudaMalloc(&dev_B, size);
  cudaMalloc(&dev C, size);
  // Copy data from host to device
  cudaMemcpy(dev A, A, size,
cudaMemcpyHostToDevice);
  cudaMemcpy(dev B, B, size,
cudaMemcpyHostToDevice);
  // Launch the kernel
  int blockSize = 256;
  int numBlocks = (n + blockSize - 1) / blockSize;
// Copy data from device to host
cudaMemcpy(C, dev C, size, cudaMemcpyDeviceToHost);
// Print the results
for (int i = 0; i < 10; i++) {
  cout << C[i] << " ";
cout << endl;
// Free memory
cudaFree(dev A);
cudaFree(dev B);
cudaFree(dev C);
cudaFreeHost(A);
cudaFreeHost(B);
cudaFreeHost(C);
return 0;
```

Matrix multiplication

Matrix-Matrix Multiplication

Consider two n x n matrices A and B partitioned into p blocks A_{ij} and B_{ij} (0
≤ i, j <) of size each.

$$\sqrt{p}$$
 $(n/\sqrt{p}) \times (n/\sqrt{p})$

- Process P_{ij} initially stores A_{ij} and B_{ij} and computes block C_{ij} of the result matrix.
- Computing submatrix C_{i,j} requires all submatrices A_{i,k} and B_{k,j} for 0 ≤ k <
- All-to-all broadcast blocks of A along rows and B along columns.
- Perform local submatrix multiplication.

Example

$$A = \begin{bmatrix} 2 & 1 & 5 & 3 \\ 0 & 7 & 1 & 6 \\ \hline 9 & 2 & 4 & 4 \\ \hline 3 & 6 & 7 & 2 \end{bmatrix}$$

B =
$$\begin{bmatrix} 6 & 1 & 2 & 3 \\ 4 & 5 & 6 & 5 \\ \hline 1 & 9 & 8 & -8 \\ 4 & 0 & -8 & 5 \end{bmatrix}$$

Example

$$A = \begin{bmatrix} 2 & 1 & 5 & 3 \\ 0 & 7 & 1 & 6 \\ 9 & 2 & 4 & 4 \\ 3 & 6 & 7 & 2 \end{bmatrix}$$

B =
$$\begin{bmatrix} 6 & 1 & 2 & 3 \\ 4 & 5 & 6 & 5 \\ 1 & 9 & 8 & -8 \\ 4 & 0 & -8 & 5 \end{bmatrix}$$

Metrics divided into 4 squares

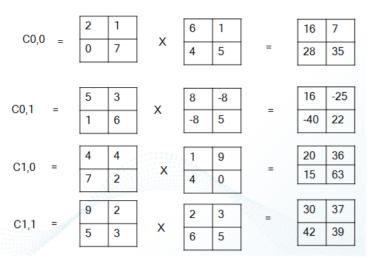
| 2 | 1 | 5 | 3 |
|---|---|---|---|
| 0 | 7 | 1 | 6 |
| 9 | 5 | 4 | 4 |
| 2 | 3 | 7 | 2 |

| 6 | 1 | | 2 | 3 |
|---|---|-----|----|----|
| 4 | 5 | | 6 | 5 |
| | | . , | | |
| 1 | 9 | | 8 | -8 |
| 4 | 0 | | -8 | 5 |

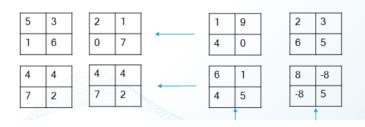
Metrics alignment

| 2 | 1 | 5 | 3 |
|---|---|---|---|
| 0 | 7 | 1 | 6 |
| 9 | 5 | 4 | 4 |
| 2 | 3 | 7 | 2 |

| 6 | 1 | 2 | 3 |
|---|---|----|----|
| 4 | 5 | 6 | 5 |
| | | | |
| 1 | 9 | 8 | -8 |
| 4 | 0 | -8 | 5 |



Now next step shifting operation is performed on A and B, Shift A one step left and B one step Up



```
#include <cuda_runtime.h> #include <iostream>
__global__void matmul(int* A, int* B, int* C, int N) { int Row =
   blockIdx.y*blockDim.y+threadIdx.y;
   int Col = blockIdx.x*blockDim.x+threadIdx.x;
   if (Row < N && Col < N) { int Pvalue = 0;
      for (int k = 0; k < N; k++) {
            Pvalue += A[Row*N+k] * B[k*N+Col];
        }
        C[Row*N+Col] = Pvalue;
   }
}
int main() {
   int N = 512;
   int size = N * N * sizeof(int);
   int* A, * B, * C;
   int* dev_A, * dev_B, * dev_C;</pre>
```

After performing local matrix multiplication

```
cudaMallocHost(&A, size);
cudaMallocHost(&B, size);
cudaMallocHost(&C, size);
cudaMalloc(&dev A, size);
```

```
cudaMalloc(&dev_C, size);

// Initialize matrices A and B
   for (int i = 0; i < N; i++) {
      for (int j = 0; j < N; j++) {
          A[i*N+j] = i*N+j;
          B[i*N+j] = j*N+i;
      }
}</pre>
```

cudaMalloc(&dev B, size);

cudaMemcpy(dev_A, A, size, cudaMemcpyHostToDevice); cudaMemcpy(dev_B, B, size, cudaMemcpyHostToDevice);

dim3 dimBlock(16, 16);
dim3 dimGrid(N/dimBlock.x, N/dimBlock.y);
matmul<<<dimGrid, dimBlock>>>(dev_A, dev_B,
dev_C, N);

cudaMemcpy(C, dev_C)

```
// Print the result
for (int i = 0; i < 10; i++) {
    for (int j = 0; j < 10; j++) {
        std::cout << C[i*N+j] << " ";
    }
    std::cout << std::endl;
}

// Free memory
cudaFree(dev_A);
cudaFree(dev_B);
cudaFree(dev_C);
cudaFreeHost(A);
cudaFreeHost(B);
cudaFreeHost(C);</pre>
```

- In this program, the `matmul` kernel takes in the two input matrices `A` and `B`, the output matrix `C`, and the size of the matrices `N`. The kernel uses the `blockIdx.x`, `blockIdx.y`, `threadIdx.x`, and `threadIdx.y` variables to calculate the indices of the current thread. If the indices are less than `N`, the kernel performs the matrix multiplication operation `Pvalue += A[Row*N+k] * B[k*N+Col]` and store the Pvalue in `C[Row*N+Col]`.
- In the `main` function, the program first allocates memory for the input and output matrices on the host and initializes them. Then it allocates memory for the matrices on the device and copies the data from the host to the device using `cudaMemcpy`.
- Next, the program launches the kernel with the appropriate grid and block dimensions. The kernel uses a 2D grid of thread blocks to perform the matrix multiplication in parallel.
- Finally, it copies the data from device to host using cudaMemcpy and prints the result using nested for loop. And it also frees the memory used.

Group 2

| Grou | p 2 |
|------|--|
| 6. | Mini Project: Evaluate performance enhancement of parallel Quicksort Algorithm using MPI |
| 7. | Mini Project: Implement Huffman Encoding on GPU |
| 8. | Mini Project: Implement Parallelization of Database Query optimization |
| 9. | Mini Project: Implement Non-Serial Polyadic Dynamic Programming with GPU Parallelization |

410251: Deep Learning

| Answers | Coding Efficiency | Viva | Viva Timely Completion | | Dated Sign of Subject Teacher | | |
|---------|----------------------|------|---------------------------|----|----------------------------------|--|--|
| 5 | 5 | 5 | 5 | 20 | | | |
| | | | | | | | |

| Start Date | Date of Completion: |
|------------|---------------------|

Group B Deep Learning

Assignment No: 1

Title of the Assignment: Linear regression by using Deep Neural network: Implement Boston housing price.prediction problem by Linear regression using Deep Neural network. Use Boston House price prediction dataset.

Objective of the Assignment: Students should be able to perform Linear regression by using Deep Neural network on Boston House Dataset.

Prerequisite:

- 1. Basic of programming language
- 2. Concept of Linear Regression
- 3. Concept of Deep Neural Network

Contents for Theory:

- 1. What is Linear Regression
- 2. Example of Linear Regression
- 3. Concept of Deep Neural Network
- 4. How Deep Neural Network Work
- 5. Code Explanation with Output

Department of Computer Engineering, PGMCOE

What is Linear Regression?

Linear regression is a statistical approach that is commonly used to model the relationship between a dependent variable and one or more independent variables. It assumes a linear relationship between the variables and uses mathematical methods to estimate the coefficients that best fit the data.

Deep neural networks are a type of machine learning algorithm that are modeled after the structure and function of the human brain. They consist of multiple layers of interconnected neurons that process data and learn from it to make predictions or classifications.

Linear regression using deep neural networks combines the principles of linear regression with the power of deep learning algorithms. In this approach, the input features are passed through one or more layers of neurons to extract features and then a linear regression model is applied to the output of the last layer to make predictions. The weights and biases of the neural network are adjusted during training to optimize the performance of the model.

This approach can be used for a variety of tasks, including predicting numerical values, such as stock prices or housing prices, and classifying data into categories, such as detecting whether an image contains a particular object or not. It is often used in fields such as finance, healthcare, and image recognition.

Example Of Linear Regression

A suitable example of linear regression using deep neural network would be predicting the price of a house based on various features such as the size of the house, the number of bedrooms, the location, and the age of the house.

In this example, the input features would be fed into a deep neural network, consisting of multiple layers of interconnected neurons. The first few layers of the network would learn to extract features from the input data, such as identifying patterns and correlations between the input features.

The output of the last layer would then be passed through a linear regression model, which would use the learned features to predict the price of the house.

During training, the weights and biases of the neural network would be adjusted to minimize the difference between the predicted price and the actual price of the house. This process is known as gradient descent, and it involves iteratively adjusting the model's parameters until the optimal values are reached.

Once the model is trained, it can be used to predict the price of a new house based on its features. This approach can be used in the real estate industry to provide accurate and reliable estimates of house prices, which can help both buyers and sellers make informed decisions.

Concept of Deep Neural Network-

A deep neural network is a type of machine learning algorithm that is modeled after the structure and function of the human brain. It consists of multiple layers of interconnected nodes, or artificial neurons, that process data and learn from it to make predictions or classifications.

Each layer of the network performs a specific type of processing on the data, such as identifying patterns or correlations between features, and passes the results to the next layer. The layers closest to the input are known as the "input layer", while the layers closest to the output are known as the "output layer".

The intermediate layers between the input and output layers are known as "hidden layers". These layers are responsible for extracting increasingly complex features from the input data, and can be deep (i.e., containing many hidden layers) or shallow (i.e., containing only a few hidden layers).

Deep neural networks are trained using a process known as backpropagation, which involves adjusting the weights and biases of the nodes based on the error between the predicted output and the actual output. This process is repeated for multiple iterations until the model reaches an optimal level of accuracy.

Deep neural networks are used in a variety of applications, such as image and speech recognition, natural language processing, and recommendation systems. They are capable of learning from vast amounts of data and can automatically extract features from raw data, making them a powerful tool for solving complex problems in a wide range of domains.

How Deep Neural Network Work-

Boston House Price Prediction is a common example used to illustrate how a deep neural network can work for regression tasks. The goal of this task is to predict the price of a house in Boston based on various features such as the number of rooms, crime rate, and accessibility to public transportation.

Here's how a deep neural network can work for Boston House Price Prediction:

- 1. **Data preprocessing:** The first step is to preprocess the data. This involves normalizing the input features to have a mean of 0 and a standard deviation of 1, which helps the network learn more efficiently. The dataset is then split into training and testing sets.
- 2. **Model architecture:** A deep neural network is then defined with multiple layers. The first layer is the input layer, which takes in the normalized features. This is followed by several hidden layers, which can be deep or shallow. The last layer is the output layer, which predicts the house price.
- 3. **Model training:** The model is then trained using the training set. During training, the weights and biases of the nodes are adjusted based on the error between the predicted output and the actual output. This is done using an optimization algorithm such as stochastic gradient descent.
- 4. Model evaluation: Once the model is trained, it is evaluated using the testing set. The

Department of Computer Engineering, PGMCOE

410255: Laboratory Practice V

performance of the model is measured using metrics such as mean squared error or mean absolute error.

- 5. **Model prediction:** Finally, the trained model can be used to make predictions on new data, such as predicting the price of a new house in Boston based on its features.
- 6. By using a deep neural network for Boston House Price Prediction, we can obtain accurate predictions based on a large set of input features. This approach is scalable and can be used for other regression tasks as well.

Boston House Price Prediction Dataset-

Boston House Price Prediction is a well-known dataset in machine learning and is often used to demonstrate regression analysis techniques. The dataset contains information about 506 houses in Boston, Massachusetts, USA. The goal is to predict the median value of owner-occupied homes in thousands of dollars.

The dataset includes 13 input features, which are:

CRIM: per capita crime rate by town

ZN: proportion of residential land zoned for lots over 25,000 sq.ft.

INDUS: proportion of non-retail business acres per town

CHAS: Charles River dummy variable (1 if tract bounds river; 0 otherwise)

NOX: nitric oxides concentration (parts per 10 million)

RM: average number of rooms per dwelling

AGE: proportion of owner-occupied units built prior to 1940

DIS: weighted distances to five Boston employment centers

RAD: index of accessibility to radial highways

TAX: full-value property-tax rate per \$10,000

PTRATIO: pupil-teacher ratio by town

B: 1000(Bk - 0.63)^2 where Bk is the proportion of black people by town

LSTAT: % lower status of the population

The output variable is the median value of owner-occupied homes in thousands of dollars (MEDV).

To predict the median value of owner-occupied homes, a regression model is trained on the dataset. The model can be a simple linear regression model or a more complex model, such as a deep neural network.

After the model is trained, it can be used to predict the median value of owner-occupied homes based on the input features. The model's accuracy can be evaluated using metrics such as mean squared error or

Department of Computer Engineering, PGMCOE

mean absolute error.

Boston House Price Prediction is a example of regression analysis and is often used to teach machine learning concepts. The dataset is also used in research to compare the performance of different regression models.

Source Code with Explanation-

#Importing the pandas for data processing and numpy for numerical computing

import numpy as np
import pandas as pd

Importing the Boston Housing dataset from the sklearn

from sklearn.datasets import load boston

boston = load boston()

#Converting the data into pandas dataframe

data = pd.DataFrame(boston.data)

#First look at the data

data.head()

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 |
|---|---------|------|------|-----|-------|-------|------|--------|-----|-------|------|--------|------|
| 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3 | 396.90 | 4.98 |
| 1 | 0.02731 | 0.0 | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8 | 396.90 | 9.14 |
| 2 | 0.02729 | 0.0 | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8 | 392.83 | 4.03 |
| 3 | 0.03237 | 0.0 | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7 | 394.63 | 2.94 |
| 4 | 0.06905 | 0.0 | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7 | 396.90 | 5.33 |

#Adding the feature names to the dataframe

data.columns = boston.feature names

#Adding the target variable to the dataset

data['PRICE'] = boston.target

#Looking at the data with names and target variable

data.head(n=10)

Department of Computer Engineering,

PGMCOE

| | | | | | | | | | | | | | 410255: La | aboratory l |
|----|---------|-------|--------|-------|-------|--------|--------|--------|-----|-------|---------|--------|------------|-------------|
| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO | В | LSTAT | PRICE |
| 0 | 0.00632 | 18.0 | 2.31 | 0.0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3 | 396.90 | 4.98 | 24.0 |
| 1 | 0.02731 | 0.0 | 7.07 | 0.0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8 | 396.90 | 9.14 | 21.6 |
| 2 | 0.02729 | 0.0 | 7.07 | 0.0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8 | 392.83 | 4.03 | 34.7 |
| 3 | 0.03237 | 0.0 | 2.18 | 0.0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7 | 394.63 | 2.94 | 33.4 |
| 4 | 0.06905 | 0.0 | 2.18 | 0.0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7 | 396.90 | 5.33 | 36.2 |
| 5 | 0.02985 | 0.0 | 2.18 | 0.0 | 0.458 | 6.430 | 58.7 | 6.0622 | 3.0 | 222.0 | 18.7 | 394.12 | 5.21 | 28.7 |
| 6 | 0.08829 | 12.5 | 7.87 | 0.0 | 0.524 | 6.012 | 66.6 | 5.5605 | 5.0 | 311.0 | 15.2 | 395.60 | 12.43 | 22.9 |
| 7 | 0.14455 | 12.5 | 7.87 | 0.0 | 0.524 | 6.172 | 96.1 | 5.9505 | 5.0 | 311.0 | 15.2 | 396.90 | 19.15 | 27.1 |
| 8 | 0.21124 | 12.5 | 7.87 | 0.0 | 0.524 | 5.631 | 100.0 | 6.0821 | 5.0 | 311.0 | 15.2 | 386.63 | 29.93 | 16.5 |
| 9 | 0.17004 | 12.5 | 7.87 | 0.0 | 0.524 | 6.004 | 85.9 | 6.5921 | 5.0 | 311.0 | 15.2 | 386.71 | 17.10 | 18.9 |
| #S | hape of | the | data | | | | | | | | | | | |
| pr | int(dat | a.sha | ape) | | | | | | | | | | | |
| #C | hecking | the | null | value | s in | the da | ataset | | | | | | | |
| da | ta.isnu | 11() | .sum() | | | | | | | | | | | |

CRIM 0 ZN 0 INDUS 0 CHAS 0 NOX RM 0 AGE 0 DIS RAD 0 TAX PTRATIO В LSTAT PRICE dtype: int64

#Checking the statistics of the data

data.describe()

This is sometimes very useful, for example if you look at the CRIM the max is 88.97 and 75% of the value is below 3.677083 and

mean is 3.613524 so it means the max values is actually an outlier or there are

Department of Computer Engineering,

PGMCOE

outliers present in the column

| | CRIM | ZN | INDUS | CHAS | NOX | RM | AGE | DIS | RAD | TAX | PTRATIO |
|-------|------------|------------|------------|------------|------------|------------|------------|------------|------------|------------|------------|
| count | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 |
| mean | 3.613524 | 11.363636 | 11.136779 | 0.069170 | 0.554695 | 6.284634 | 68.574901 | 3.795043 | 9.549407 | 408.237154 | 18.455534 |
| std | 8.601545 | 23.322453 | 6.860353 | 0.253994 | 0.115878 | 0.702617 | 28.148861 | 2.105710 | 8.707259 | 168.537116 | 2.164946 |
| min | 0.006320 | 0.000000 | 0.460000 | 0.000000 | 0.385000 | 3.561000 | 2.900000 | 1.129600 | 1.000000 | 187.000000 | 12.600000 |
| 25% | 0.082045 | 0.000000 | 5.190000 | 0.000000 | 0.449000 | 5.885500 | 45.025000 | 2.100175 | 4.000000 | 279.000000 | 17.400000 |
| 50% | 0.256510 | 0.000000 | 9.690000 | 0.000000 | 0.538000 | 6.208500 | 77.500000 | 3.207450 | 5.000000 | 330.000000 | 19.050000 |
| 75% | 3.677083 | 12.500000 | 18.100000 | 0.000000 | 0.624000 | 6.623500 | 94.075000 | 5.188425 | 24.000000 | 666.000000 | 20.200000 |
| max | 88.976200 | 100.000000 | 27.740000 | 1.000000 | 0.871000 | 8.780000 | 100.000000 | 12.126500 | 24.000000 | 711.000000 | 22.000000 |

data.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 506 entries, 0 to 505

Data columns (total 14 columns):

| # | Column | Non-Null Count | Dtype |
|----|---------|----------------|---------|
| | | | |
| 0 | CRIM | 506 non-null | float64 |
| 1 | ZN | 506 non-null | float64 |
| 2 | INDUS | 506 non-null | float64 |
| 3 | CHAS | 506 non-null | float64 |
| 4 | NOX | 506 non-null | float64 |
| 5 | RM | 506 non-null | float64 |
| 6 | AGE | 506 non-null | float64 |
| 7 | DIS | 506 non-null | float64 |
| 8 | RAD | 506 non-null | float64 |
| 9 | TAX | 506 non-null | float64 |
| 10 | PTRATIO | 506 non-null | float64 |
| 11 | В | 506 non-null | float64 |
| 12 | LSTAT | 506 non-null | float64 |
| 13 | PRICE | 506 non-null | float64 |

dtypes: float64(14)
memory usage: 55.5 KB

Department of Computer Engineering,

#checking the distribution of the target variable

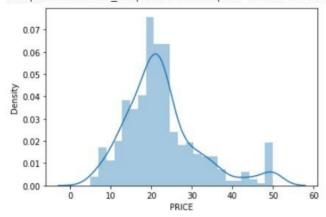
import seaborn as sns

sns.distplot(data.PRICE)

#The distribution seems normal, has not be the data normal we would have perform log transformation or took to square root of the data to make the data normal.

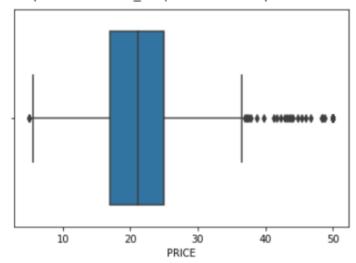
Normal distribution is need for the machine learning for better predictiblity of the model

<matplotlib.axes. subplots.AxesSubplot at 0x7f44d082c670>



#Distribution using box plot
sns.boxplot(data.PRICE)

<matplotlib.axes._subplots.AxesSubplot at 0x7f44d077ed60>



#Checking the correlation of the independent feature with the dependent feature # Correlation is a statistical technique that can show whether and how strongly pairs of variables are related. An intelligent correlation analysis can lead to a greater understanding of your data

#checking Correlation of the data

correlation = data.corr()

correlation.loc['PRICE']

CRIM -0.388305

Department of Computer Engineering,

```
ZN
           0.360445
INDUS
          -0.483725
           0.175260
CHAS
NOX
          -0.427321
           0.695360
RM
          -0.376955
AGE
DIS
           0.249929
RAD
          -0.381626
TAX
          -0.468536
PTRATIO
          -0.507787
           0.333461
В
          -0.737663
LSTAT
PRICE
           1.000000
```

Name: PRICE, dtype: float64

plotting the heatmap

import matplotlib.pyplot as plt

fig,axes = plt.subplots(figsize=(15,12))

sns.heatmap(correlation, square = True, annot = True)

By looking at the correlation plot LSAT is negatively correlated with -0.75 and RM is positively correlated to the price and PTRATIO is correlated negatively with -0.51



Department of Computer Engineering,

```
# Checking the scatter plot with the most correlated features
 plt.figure(figsize = (20,5))
 features = ['LSTAT','RM','PTRATIO']
 for i, col in enumerate(features):
     plt.subplot(1, len(features), i+1)
     x = data[col]
     y = data.PRICE
     plt.scatter(x, y, marker='o')
     plt.title("Variation in House prices")
     plt.xlabel(col)
     plt.ylabel('"House prices in $1000"')
            Variation in House prices
                                 prices in $1000"
                                                                prices in $1000
   20
                                  20
                                                                 20
                                                                              PTRATIO
 # Splitting the dependent feature and independent feature
 #X = data[['LSTAT','RM','PTRATIO']]
 X = data.iloc[:,:-1]
 y= data.PRICE
 # In order to provide a standardized input to our neural network, we need the
 perform the normalization of our dataset.
 # This can be seen as an step to reduce the differences in scale that may arise
 from the existent features.
 # We perform this normalization by subtracting the mean from our data and
 dividing it by the standard deviation.
 # One more time, this normalization should only be performed by using the mean
 and standard deviation from the training set,
 # in order to avoid any information leak from the test set.
 mean = X train.mean(axis=0)
 std = X train.std(axis=0)
 X train = (X train - mean) / std
 X \text{ test} = (X \text{ test - mean}) / \text{std}
 #Linear Regression
 from sklearn.linear model import LinearRegression
Department of Computer Engineering,
PGMCOE
```

```
regressor = LinearRegression()
#Fitting the model
regressor.fit(X train,y train)
# Model Evaluation
#Prediction on the test dataset
y pred = regressor.predict(X test)
# Predicting RMSE the Test set results
from sklearn.metrics import mean squared error
rmse = (np.sqrt(mean squared error(y test, y pred)))
print(rmse)
from sklearn.metrics import r2 score
r2 = r2 score(y test, y pred)
print(r2)
# Neural Networks
#Scaling the dataset
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X train = sc.fit transform(X train)
X test = sc.transform(X test)
# Due to the small amount of presented data in this dataset, we must be careful
to not create an overly complex model,
# which could lead to overfitting our data. For this, we are going to adopt an
architecture based on two Dense layers,
# the first with 128 and the second with 64 neurons, both using a ReLU activation
function.
# A dense layer with a linear activation will be used as output layer.
# In order to allow us to know if our model is properly learning, we will use a
mean squared error loss function and to report the performance of it we will
adopt the mean average error metric.
# By using the summary method from Keras, we can see that we have a total of
10,113 parameters, which is acceptable for us.
#Creating the neural network model
import keras
from keras.layers import Dense, Activation, Dropout
from keras.models import Sequential
model = Sequential()
model.add(Dense(128,activation = 'relu',input dim =13))
model.add(Dense(64,activation = 'relu'))
```

Department of Computer Engineering,

```
model.add(Dense(32,activation = 'relu'))
model.add(Dense(16,activation = 'relu'))
model.add(Dense(1))
#model.compile(optimizer='adam', loss='mse', metrics=['mae'])
model.compile(optimizer = 'adam',loss = 'mean squared error',metrics=['mae'])
!pip install ann visualizer
!pip install graphviz
from ann visualizer.visualize import ann viz;
#Build your model here
ann viz(model, title="DEMO ANN");
history = model.fit(X train, y train, epochs=100, validation split=0.05)
# By plotting both loss and mean average error, we can see that our model was
capable of learning patterns in our data without overfitting taking place (as
shown by the validation set curves)
from plotly.subplots import make subplots
import plotly.graph objects as go
fig = go.Figure()
fig.add trace(go.Scattergl(y=history.history['loss'],
                    name='Train'))
fig.add trace(go.Scattergl(y=history.history['val loss'],
                    name='Valid'))
fig.update layout (height=500, width=700,
                  xaxis title='Epoch',
                  yaxis title='Loss')
fig.show()
                                                     - Valid
   500
   400
   300
   200
   100
    0
             20
                         Epoch
```

Department of Computer Engineering,

```
fig.add trace(go.Scattergl(y=history.history['mae'],
                       name='Train'))
 fig.add trace(go.Scattergl(y=history.history['val mae'],
                       name='Valid'))
 fig.update layout (height=500, width=700,
                    xaxis title='Epoch',
                     yaxis title='Mean Absolute Error')
 fig.show()
                                                                 - Train
                                                                Valid
     20
 Mean Absolute Error
     15
     10
      5
      0
                 20
                                                 80
                               Epoch
 #Evaluation of the model
 y pred = model.predict(X test)
 mse nn, mae nn = model.evaluate(X test, y test)
 print('Mean squared error on test data: ', mse nn)
 print('Mean absolute error on test data: ', mae nn)
 4/4 [==============] - 0s 4ms/step - loss: 10.5717 - mae: 2.2670
 Mean squared error on test data: 10.571733474731445
 Mean absolute error on test data: 2.2669904232025146
 #Comparison with traditional approaches
 #First let's try with a simple algorithm, the Linear Regression:
 from sklearn.metrics import mean absolute error
 lr model = LinearRegression()
 lr_model.fit(X_train, y_train)
Department of Computer Engineering,
```

```
y pred lr = lr model.predict(X test)
mse lr = mean squared error(y test, y pred lr)
mae lr = mean absolute error(y test, y pred lr)
print('Mean squared error on test data: ', mse lr)
print('Mean absolute error on test data: ', mae lr)
from sklearn.metrics import r2 score
r2 = r2 \text{ score}(y \text{ test, } y \text{ pred})
print(r2)
0.8812832788381159
 # Predicting RMSE the Test set results
 from sklearn.metrics import mean squared error
 rmse = (np.sqrt(mean squared error(y test, y pred)))
 print(rmse)
 3.320768607496587
 # Make predictions on new data
 import sklearn
 new data = sklearn.preprocessing.StandardScaler().fit transform(([[0.1, 10.0,
 5.0, 0, 0.4, 6.0, 50, 6.0, 1, 400, 20, 300, 10]]))
 prediction = model.predict(new data)
 print("Predicted house price:", prediction)
 1/1 [======= ] - 0s 70ms/step
 Predicted house price: [[11.104753]]
 #new data
 sklearn.preprocessing.StandardScaler().fit transform(([[0.1,
 5.0, 0, 0.4, 6.0, 50, 6.0, 1, 400, 20, 300, 10]])) is a line of code
 that standardizes the input features of a new data point.
 In this specific case, we have a new data point represented as a
 list of 13 numeric values ([0.1, 10.0, 5.0, 0, 0.4, 6.0, 50, 6.0, 1,
 400, 20, 300, 10]) that represents the values for the 13 features of
 the Boston House Price dataset.
 The StandardScaler() function from the sklearn.preprocessing module is
 used to standardize the data. Standardization scales each featureto
 have zero mean and unit variance, which is a common preprocessingstep
 in machine learning to ensure that all features contribute equally to
 the model.
 The fit transform() method is used to fit the scaler to the data and
 apply the standardization transformation. The result is a new data
```

point with standardized feature values.

Conclusion- In this way we can Predict the Boston House Price using Deep Neural Network.

Assignment Question

- 1. What is Linear Regression?
- 2. What is a Deep Neural Network?
- 3. What is the concept of standardization?
- 4. Why split data into train and test?
- 5. Write Down Application of Deep Neural Network?

Department of Computer Engineering,

| Answers | Coding Efficiency | Viva | Timely Completion | Total | Dated Sign of Subject Teacher |
|---------|----------------------|------|----------------------|-------|----------------------------------|
| 5 | 5 | 5 | 5 | 20 | |
| | | | | | |

| Start Date | • | Date of Completion: |
|------------|---|---------------------|
| Start Date | | Date of Completion |

Assignment No: 2A

Title of the Assignment: Binary classification using Deep Neural Networks Example: Classify movie reviews into positive" reviews and "negative" reviews, just based on the text content of the reviews. Use IMDB dataset

Objective of the Assignment: Students should be able to Classify movie reviews into positive reviews and "negative reviews on IMDB Dataset.

Prerequisite:

- 1. Basic of programming language
- 2. Concept of Classification
- Concept of Deep Neural Network

Contents for Theory:

- 1. What is Classification
- 2. Example of Classification
- 3. How Deep Neural Network Work on Classification
- 4. Code Explanation with Output

What is Classification?

Classification is a type of supervised learning in machine learning that involves categorizing data into predefined classes or categories based on a set of features or characteristics. It is used to predict the classof new, unseen data based on the patterns learned from the labeled training data.

In classification, a model is trained on a labeled dataset, where each data point has a known class label. The model learns to associate the input features with the corresponding class labels and can then be used to classify new, unseen data.

For example, we can use classification to identify whether an email is spam or not based on its content and metadata, to predict whether a patient has a disease based on their medical records and symptoms, orto classify images into different categories based on their visual features.

Classification algorithms can vary in complexity, ranging from simple models such as decision trees andknearest neighbors to more complex models such as support vector machines and neural networks. The choice of algorithm depends on the nature of the data, the size of the dataset, and the desired level of accuracy and interpretability.

Classification is a common task in deep neural networks, where the goal is to predict the class of aninput based on its features. Here's an example of how classification can be performed in a deep neural network using the popular MNIST dataset of handwritten digits.

The MNIST dataset contains 60,000 training images and 10,000 testing images of handwritten digits from 0 to 9. Each image is a grayscale 28x28 pixel image, and the task is to classify each image into one of the 10 classes corresponding to the 10 digits.

We can use a convolutional neural network (CNN) to classify the MNIST dataset. A CNN is a type of deep neural network that is commonly used for image classification tasks.

How Deep Neural Network Work on Classification-

Deep neural networks are commonly used for classification tasks because they can automatically learn to extract relevant features from raw input data and map them to the correct output class.

The basic architecture of a deep neural network for classification consists of three main parts: an input layer, one or more hidden layers, and an output layer. The input layer receives the raw input data, which is usually preprocessed to a fixed size and format. The hidden layers are composed of neurons that apply linear transformations and nonlinear activations to the input features to extract relevant patterns and representations. Finally, the output layer produces the predicted class labels, usually as a probability distribution over the possible classes.

During training, the deep neural network learns to adjust its weights and biases in each layer to minimize the difference between the predicted output and the true labels. This is typically done by optimizing a loss function that measures the discrepancy between the predicted and true labels, using techniques such as gradient descent or stochastic gradient descent.

One of the key advantages of deep neural networks for classification is their ability to learn hierarchical representations of the input data. In a deep neural network with multiple hidden layers, each layer learns to capture more complex and abstract features than the previous layer, by building on the representations learned by the earlier layers. This hierarchical structure allows deep neural networks to learn highly discriminative features that can separate different classes of input data, even when the data is highly complex or noisy.

Overall, the effectiveness of deep neural networks for classification depends on the choice of architecture, hyperparameters, and training procedure, as well as the quality and quantity of the training data. When trained properly, deep neural networks can achieve state-of-the-art performance on a wide range of classification tasks, from image recognition to natural language processing.

IMDB Dataset-The IMDB dataset is a large collection of movie reviews collected from the IMDB website, which is a popular source of user-generated movie ratings and reviews. The dataset consists of 50,000 movie reviews, split into 25,000 reviews for training and 25,000 reviews for testing.

Each review is represented as a sequence of words, where each word is represented by an integer index based on its frequency in the dataset. The labels for each review are binary, with 0 indicating a negative review and 1 indicating a positive review.

The IMDB dataset is commonly used as a benchmark for sentiment analysis and text classification tasks, where the goal is to classify the movie reviews as either positive or negative based on their text content. The dataset is challenging because the reviews are often highly subjective and can contain complex language and nuances of meaning, making it difficult for traditional machine learning approaches to accurately classify them.

Deep learning approaches, such as deep neural networks, have achieved state-of-the-art performance on the IMDB dataset by automatically learning to extract relevant features from the raw text data and map them to the correct output class. The IMDB dataset is widely used in research and education for natural language processing and machine learning, as it provides a rich source of labeled text data for training and testing deep learning models.

Source Code and Output-

- # The IMDB sentiment classification dataset consists of 50,000 movie reviews from IMDB users that are labeled as either positive (1) or negative (0).
- # The reviews are preprocessed and each one is encoded as a sequence of word indexes in the form of integers.
- # The words within the reviews are indexed by their overall frequency within the dataset. For example, the integer "2" encodes the second most frequent word in the data.
- # The 50,000 reviews are split into 25,000 for training and 25,000 for testing.
- # Text Process word by word at diffrent timestamp (You may use RNN LSTM GRU)# convert input text to vector reprent input text
- # DOMAIN: Digital content and entertainment industry
- # CONTEXT: The objective of this project is to build a text classification model that analyses the customer's sentiments based on their reviews in the IMDB database. The model uses a complex deep learning model to build an embedding layer followed by a classification algorithm to analyse the sentiment of the customers.
- # DATA DESCRIPTION: The Dataset of 50,000 movie reviews from IMDB, labelled by sentiment (positive/negative).
- # Reviews have been preprocessed, and each review is encoded as a sequence of word indexes (integers).
- # For convenience, the words are indexed by their frequency in the dataset, meaning the for that has index 1 is the most frequent word.
- # Use the first 20 words from each review to speed up training, using a max vocabulary size of 10,000.
- # As a convention, "0" does not stand for a specific word, but instead is used to encode any unknownword.
- # PROJECT OBJECTIVE: Build a sequential NLP classifier which can use input text parameters to determine the customer sentiments.

import numpy as np import pandas as pd

from sklearn.model_selection import train_test_split

#loading imdb data with most frequent 10000 words

```
from keras.datasets import imdb
(X train, y train), (X test, y test) = imdb.load data(num words=10000) # you may take top 10,000word
frequently used review of movies other are discarded
#consolidating data for EDA Exploratory data analysis (EDA) is used by data scientists to analyze andinvestigate
data sets and summarize their main characteristics
data = np.concatenate((X train, X test), axis=0) # axis 0 is first running vertically downwards acrossrows
(axis 0), axis 1 is second running horizontally across columns (axis 1),
label = np.concatenate((y_train, y_test), axis=0)
X_train.shape
(25000,)
X test.shape
(25000,)
y_train.shape
(25000,)
y_test.shape
(25000,)
print("Review is ",X train[0]) # series of no converted word to vocabulory associated with index
print("Review is ",y train[0])
Review is [1, 194, 1153, 194, 8255, 78, 228, 5, 6, 1463, 4369, 5012, 134, 26, 4, 715, 8, 118, 1634, 14,
394, 20, 13, 119, 954, 189, 102, 5, 207, 110, 3103, 21, 14, 69, 188, 8, 30, 23, 7, 4, 249, 126, 93, 4, 114,
9, 2300, 1523, 5, 647, 4, 116, 9, 35, 8163, 4, 229, 9, 340, 1322, 4, 118, 9, 4, 130, 4901, 19, 4, 1002, 5,
89, 29, 952, 46, 37, 4, 455, 9, 45, 43, 38, 1543, 1905, 398, 4, 1649, 26, 6853, 5, 163, 11, 3215, 2, 4,
1153, 9, 194, 775, 7, 8255, 2, 349, 2637, 148, 605, 2, 8003, 15, 123, 125, 68, 2, 6853, 15, 349, 165,
4362, 98, 5, 4, 228, 9, 43, 2, 1157, 15, 299, 120, 5, 120, 174, 11, 220, 175, 136, 50, 9, 4373, 228, 8255,
5, 2, 656, 245, 2350, 5, 4, 9837, 131, 152, 491, 18, 2, 32, 7464, 1212, 14, 9, 6, 371, 78, 22, 625, 64,
1382, 9, 8, 168, 145, 23, 4, 1690, 15, 16, 4, 1355, 5, 28, 6, 52, 154, 462, 33, 89, 78, 285, 16, 145, 95]
Review is 0
vocab=imdb.get_word_index() # Retrieve the word index file mapping words to indices
print(vocab)
{'fawn': 34701, 'tsukino': 52006, 'nunnery': 52007, 'sonja': 16816, 'vani': 63951, 'woods': 1408, 'spiders':
16115,
y_train
array([1, 0, 0, ..., 0, 1, 0])
y_test
array([0, 1, 1, ..., 0, 0, 0])
# Function to perform relevant sequence adding on the data
# Now it is time to prepare our data. We will vectorize every review and fill it with zeros so that it
```

```
contains exactly 10000 numbers.
# That means we fill every review that is shorter than 500 with zeros.
# We do this because the biggest review is nearly that long and every input for our neural network needsto
have the same size.
# We also transform the targets into floats.
# sequences is name of method the review less than 10000 we perform padding overthere#
binary vectorization code:
# VECTORIZE as one cannot feed integers into a NN
# Encoding the integer sequences into a binary matrix - one hot encoder basically
# From integers representing words, at various lengths - to a normalized one hot encoded tensor (matrix) of
10k columns
def vectorize(sequences, dimension = 10000):
# We will vectorize every review and fill it with zerosso that it contains exactly 10,000 numbers.
  # Create an all-zero matrix of shape (len(sequences), dimension)
   results = np.zeros((len(sequences), dimension))
   for i, sequence in enumerate(sequences):
      results[i, sequence] = 1
   return results
# Now we split our data into a training and a testing set.
# The training set will contain reviews and the testing set
## Set a VALIDATION set
test_x = data[:10000]
test_y = label[:10000]
train_x = data[10000:]
train_y = label[10000:]
test_x.shape
(10000,)
test_y.shape
(10000,)
train_x.shape
(40000,)
train_y.shape
(40000,)
print("Categories:", np.unique(label))
print("Number of unique words:", len(np.unique(np.hstack(data))))
 Department of Computer Engineering,
 PGMCOE
```

```
# The hstack() function is used to stack arrays in sequence horizontally (column wise).
Categories: [0 1]
Number of unique words: 9998
length = [len(i) for i in data]
print("Average Review length:", np.mean(length))
print("Standard Deviation:", round(np.std(length)))
# The whole dataset contains 9998 unique words and the average review length is 234 words, with a
standard deviation of 173 words.
Average Review length: 234.75892
Standard Deviation: 173
# If you look at the data you will realize it has been already pre-processed.
# All words have been mapped to integers and the integers represent the words sorted by their frequency.#
This is very common in text analysis to represent a dataset like this.
# So 4 represents the 4th most used word,# 5
the 5th most used word and so on...
# The integer 1 is reserved for the start marker,
# the integer 2 for an unknown word and 0 for padding.#
Let's look at a single training example:
print("Label:", label[0])
Label: 1
print("Label:", label[1])
Label: 0
print(data[0])
# Retrieves a dict mapping words to their index in the IMDB dataset.
index = imdb.get word index() # word to index
# Create inverted index from a dictionary with document ids as keys and a list of terms as values foreach
document
reverse_index = dict([(value, key) for (key, value) in index.items()]) # id to word
decoded = " ".join( [reverse_index.get(i - 3, "#") for i in data[0]] )
# The indices are offset by 3 because 0, 1 and 2 are reserved indices for "padding", "start of sequence" and
"unknown".
print(decoded)
# this film was just brilliant casting location scenery story direction everyone's really suited the part they
 Department of Computer Engineering,
 PGMCOE
```

played and you could just imagine being there robert # is an amazing actor and now the same being director # father came from the same scottish island as myself so i loved the fact there was a real connection with this film the witty remarks throughout the film

#Adding sequence to data

Vectorization is the process of converting textual data into numerical vectors and is a process that is usually applied once the text is cleaned.

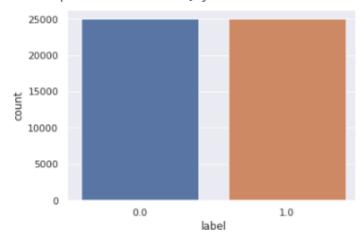
data = vectorize(data)

label = np.array(label).astype("float32")

labelDF=pd.DataFrame({'label':label})

sns.countplot(x='label', data=labelDF)

<AxesSubplot:xlabel='label', ylabel='count'>



Creating train and test data set

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(data,label, test_size=0.20, random_state=1)

X train.shape (40000,

10000)

X_test.shape (10000,

10000)

Let's create sequential model

from keras.utils import to_categorical

from keras import models

from keras import layers model

= models.Sequential()# Input -

Layer

Note that we set the input-shape to 10,000 at the input-layer because our reviews are 10,000 integerslong.

The input-layer takes 10,000 as input and outputs it with a shape of 50.

model.add(layers.Dense(50, activation = "relu", input_shape=(10000,)))

Department of Computer Engineering,

```
# Hidden - Layers
```

Please note you should always use a dropout rate between 20% and 50%. # here in our case 0.3 means 30% dropout we are using dropout to prevent overfitting.

By the way, if you want you can build a sentiment analysis without LSTMs, then you simply need to replace it by a flatten layer:

```
model.add(layers.Dropout(0.3, noise_shape=None, seed=None))
```

model.add(layers.Dense(50, activation = "relu"))

model.add(layers.Dropout(0.2, noise_shape=None, seed=None))

model.add(layers.Dense(50, activation = "relu"))

Output- Layer

model.add(layers.Dense(1, activation = "sigmoid"))

model.summary()

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|---------------------|--------------|---------|
| dense (Dense) | (None, 50) | 500050 |
| dropout (Dropout) | (None, 50) | 0 |
| dense_1 (Dense) | (None, 50) | 2550 |
| dropout_1 (Dropout) | (None, 50) | 0 |
| dense_2 (Dense) | (None, 50) | 2550 |
| dense_3 (Dense) | (None, 1) | 51 |

```
Total params: 505,201
Trainable params: 505,201
Non-trainable params: 0
```

#For early stopping

Stop training when a monitored metric has stopped improving.

monitor: Quantity to be monitored.

patience: Number of epochs with no improvement after which training will be stopped.

import tensorflow as tf

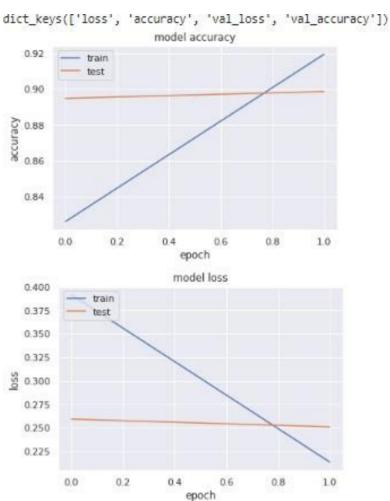
callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=3)

We use the "adam" optimizer, an algorithm that changes the weights and biases during training.

We also choose binary-crossentropy as loss (because we deal with binary

```
classification) and accuracy as our evaluation metric.
model.compile(
optimizer = "adam",
loss = "binary crossentropy",
metrics = ["accuracy"]
from sklearn.model selection import train test split
results = model.fit(
X train, y train,
epochs= 2,
batch size = 500,
validation data = (X test, y test),
callbacks=[callback]
# Let's check mean accuracy of our model
print(np.mean(results.history["val accuracy"]))
# Evaluate the model
score = model.evaluate(X test, y test, batch size=500)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
20/20 [============== ] - 1s 24ms/step - loss: 0.2511 - accuracy:
0.8986
Test loss: 0.25108325481414795
Test accuracy: 0.8985999822616577
#Let's plot training history of our model.
# list all data in history
print(results.history.keys())
# summarize history for accuracy
plt.plot(results.history['accuracy'])
plt.plot(results.history['val accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
# summarize history for loss
plt.plot(results.history['loss'])
plt.plot(results.history['val loss'])
plt.title('model loss')
plt.ylabel('loss')
Department of Computer Engineering, PGMCOE
```

```
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



Conclusion- In this way we can Classify the Movie Reviews by using DNN.

Assignment Question

- 1. What is Binary Classification?
- 2. What is binary Cross Entropy?
- 3. What is Validation Split?
- 4. What is the Epoch Cycle?
- 5. What is Adam Optimizer?

| Answers | Coding Efficiency | Viva | Timely Completion | Total | Dated Sign of Subject Teacher |
|---------|----------------------|------|----------------------|-------|----------------------------------|
| 5 | 5 | 5 | 5 | 20 | |
| | | | | | |

| Start Date: | Date of Completion: |
|-------------|---------------------|
| | |

Assignment No: 2B

Title of the Assignment: Multiclass classification using Deep Neural Networks: Example: Use the OCRletter recognition dataset https://archive.ics.uci.edu/ml/datasets/letter+recognition

Objective of the Assignment: Students should be able to solve Multiclass classification using Deep Neural NetworksSolve

Prerequisite:

- 1. Basic of programming language
- 2. Concept of Multi Classification
- 3. Concept of Deep Neural Network

Contents for Theory:

- 1. What is Multi-Classification
- 2. Example of Multi-Classification
- 3. How Deep Neural Network Work on Multi-Classification
- 4. Code Explanation with Output

Department of Computer Engineering,

What is multiclass classification?

Multi Classification, also known as multiclass classification or multiclass classification problem, is a type of classification problem where the goal is to assign input data to one of three or more classes or categories. In other words, instead of binary classification, where the goal is to assign input data to one of two classes (e.g., positive or negative), multiclass classification involves assigning input data to one of several possible classes or categories (e.g., animal species, types of products, etc.).

In multiclass classification, each input sample is associated with a single class label, and the goal of the model is to learn a function that can accurately predict the correct class label for new, unseen input data. Multiclass classification can be approached using a variety of machine learning algorithms, including decision trees, support vector machines, and deep neural networks.

Some examples of multiclass classification problems include image classification, where the goal is to classify images into one of several categories (e.g., animals, vehicles, buildings), and text classification, where the goal is to classify text documents into one of several categories (e.g., news topics, sentiment analysis).

Example of multiclass classification-

Here are a few examples of multiclass classification problems:

Image classification: The goal is to classify images into one of several categories. For example, an image classification model might be trained to classify images of animals into categories such as cats, dogs, and birds.

Text classification: The goal is to classify text documents into one of several categories. For example, a text classification model might be trained to classify news articles into categories such as politics, sports, and entertainment.

Disease diagnosis: The goal is to diagnose patients with one of several diseases based on their symptoms and medical history. For example, a disease diagnosis model might be trained to classify patients into categories such as diabetes, cancer, and heart disease.

Speech recognition: The goal is to transcribe spoken words into text. A speech recognition model might be trained to recognize spoken words in several languages or dialects.

Credit risk analysis: The goal is to classify loan applicants into categories such as low risk, medium risk, and high risk. A credit risk analysis model might be trained to classify loan applicants based on their creditscore, income, and other factors.

In all of these examples, the goal is to assign input data to one of several possible classes or categories.

Department of Computer Engineering,

Multiclass classification is a common task in machine learning and can be approached using a variety of algorithms, including decision trees, support vector machines, and deep neural networks.

Source Code and Output

import numpy as np
from tensorflow.keras.models import Sequential from
tensorflow.keras.layers import Dense, Dropoutfrom
tensorflow.keras.optimizers import RMSprop from
tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt
from sklearn import metrics
Load the OCR dataset

- # The MNIST dataset is a built-in dataset provided by Keras.
- # It consists of 70,000 28x28 grayscale images, each of which displays a single handwritten digit from 0to 9.
- # The training set consists of 60,000 images, while the test set has 10,000 images.

(x_train, y_train), (x_test, y_test) = mnist.load_data()

X_train and X_test are our array of images while y_train and y_test are our array of labels for eachimage. # The first tuple contains the training set features (X_train) and the training set labels (y_train).# The second tuple contains the testing set features (X_test) and the testing set labels (y_test). # For example, if the image shows a handwritten 7, then the label will be the intger 7.

plt.imshow(x_train[0], cmap='gray') # imshow() function which simply displays an image.
plt.show() # cmap is responsible for mapping a specific colormap to the values found in the array thatyou passed as the first argument.

- # This is because of the format that all the images in the dataset have:
- # 1. All the images are grayscale, meaning they only contain black, white and grey.# 2.

The images are 28 pixels by 25 pixels in size (28x28).

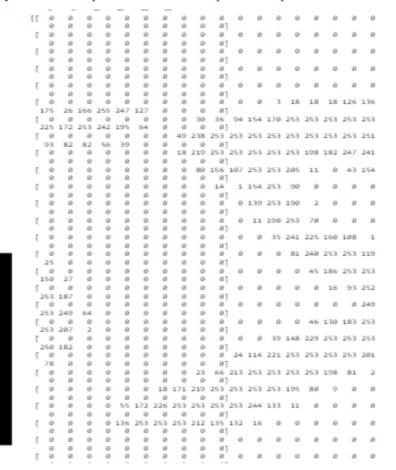
print(x_train[0])

- # image data is just an array of digits. You can almost make out a 5 from the pattern of the digits in thearray.
- # Array of 28 values

Department of Computer Engineering,

a grayscale pixel is stored as a digit between 0 and 255 where 0 is black, 255 is white and values inbetween are different shades of gray.

Therefore, each value in the [28][28] array tells the computer which color to put in that position when.



reformat our X_train array and our X_test array because they do not have the correct shape.#

25

Reshape the data to fit the model

0

10

15

20

25

print("X_train shape", x_train.shape)

10

15

20

print("y_train shape", y_train.shape)

print("X_test shape", x_test.shape)

print("y_test shape", y_test.shape)

Here you can see that for the training sets we have 60,000 elements and the testing sets have 10,000 elements.

y_train and y_test only have 1 dimensional shapes because they are just the labels of each element.

x_train and x_test have 3 dimensional shapes because they have a width and height (28x28 pixels) foreach element.

(60000, 28, 28) 1st parameter in the tuple shows us how much image we have 2nd and 3rd parameters are the pixel values from x to y (28x28)

The pixel value varies between 0 to 255.

(60000,) Training labels with integers from 0-9 with dtype of uint8. It has the shape (60000,).

Department of Computer Engineering,

```
# (10000, 28, 28) Testing data that consists of grayscale images. It has the shape (10000, 28, 28) and the dtype of
uint8. The pixel value varies between 0 to 255.
# (10000,) Testing labels that consist of integers from 0-9 with dtype uint8. It has the shape (10000,).
X_train shape (60000, 28, 28)
y train shape (60000,)
X_test shape (10000, 28, 28)
y_test shape (10000,)
# X: Training data of shape (n_samples, n_features)
# y: Training label values of shape (n_samples, n_labels)
# 2D array of height and width, 28 pixels by 28 pixels will just become 784 pixels (28 squared).
# Remember that X train has 60,000 elements, each with 784 total pixels so will become shape (60000,784).
# Whereas X test has 10,000 elements, each with each with 784 total pixels so will become shape(10000,
784).
x_{train} = x_{train.reshape}(60000, 784)
x_{test} = x_{test.reshape}(10000, 784)
x train = x train.astype('float32') # use 32-bit precision when training a neural network, so at one point the
training data will have to be converted to 32 bit floats. Since the dataset fits easily in RAM, we mightas well
convert to float immediately.
x \text{ test} = x \text{ test.astype('float32')}
x_t = 255 \# Each image has Intensity from 0 to 255x_test
/= 255
# Regarding the division by 255, this is the maximum value of a byte (the input feature's type before the conversion
to float32),
# so this will ensure that the input features are scaled between 0.0 and 1.0.#
Convert class vectors to binary class matrices
num classes = 10
y_train = np.eye(num_classes)[y_train]
                                             # Return a 2-D array with ones on the diagonal and zeros
elsewhere.
y test = np.eye(num classes)[y test] # f your particular categories is present then it mark as 1 else 0 inremain
# Define the model architecture
model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,))) # Input cosist of 784 Neuron ie 784 input,512 in
the hidden layer
model.add(Dropout(0.2)) # DROP OUT RATIO 20%
```

Department of Computer Engineering,

```
model.add(Dense(512, activation='relu')) #returns a sequence of another vectors of dimension 512
model.add(Dropout(0.2))
model.add(Dense(num classes, activation='softmax')) # 10 neurons ie output node in the output layer.#
Compile the model
 model.compile(loss='categorical crossentropy', # for a multi-class classification problemoptimizer=RMSprop(),
         metrics=['accuracy'])#
Train the model
batch size = 128 # batch size argument is passed to the layer to define a batch size for the inputs.
epochs = 20
history = model.fit(x_train, y_train, batch_size=batch_size,epochs=epochs, verbose=1,
# verbose=1 will show you an animated progress bar eg. [=======]
validation_data=(x_test, y_test))
# Using validation data means you are providing the training set and validation set yourself,
# 60000image/128=469 batch each
# Evaluate the model
score = model.evaluate(x test, y test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
Test loss: 0.08541901409626007
Test accuracy: 0.9851999878883362
```

Conclusion- In this way we can do Multi classification using DNN.

Assignment Question

- 1. What is Batch Size?
- 2. What is Dropout?
- 3. What is RMSprop?
- 4. What is the Softmax Function?
- 5. What is the Relu Function?

Department of Computer Engineering,

| Answers | Coding Efficiency | Viva | Timely Completion | Total | Dated Sign of Subject Teacher |
|---------|----------------------|------|----------------------|-------|----------------------------------|
| 5 | 5 | 5 | 5 | 20 | |
| | | | | | |

| a | | D . CO 1.1 |
|-------------|---|---------------------|
| Stort Linto | • | Date of Completion: |
| Siali Laie | | Date of Completion |
| | | |

Assignment No: 3B

Title of the Assignment: Use MNIST Fashion Dataset and create a classifier to classify fashion clothing into categories.

Objective of the Assignment: Students should be able to Classify movie reviews into positive reviews and "negative reviews on IMDB Dataset.

Prerequisite:

- 1. Basic of programming language
- 2. Concept of Classification
- 3. Concept of Deep Neural Network

Contents for Theory:

- 1. What is Classification
- 2. Example of Classification
- 3. What is CNN?
- 4. How Deep Neural Network Work on Classification
- 5. Code Explanation with Output

Department of Computer Engineering,

What is Classification?

Classification is a type of supervised learning in machine learning that involves categorizing data into predefined classes or categories based on a set of features or characteristics. It is used to predict the class of new, unseen data based on the patterns learned from the labeled training data.

In classification, a model is trained on a labeled dataset, where each data point has a known class label. The model learns to associate the input features with the corresponding class labels and can then be used to classify new, unseen data.

For example, we can use classification to identify whether an email is spam or not based on its content and metadata, to predict whether a patient has a disease based on their medical records and symptoms, or to classify images into different categories based on their visual features.

Classification algorithms can vary in complexity, ranging from simple models such as decision trees and k-nearest neighbors to more complex models such as support vector machines and neural networks. The choice of algorithm depends on the nature of the data, the size of the dataset, and the desired level of accuracy and interpretability.

Example- Classification is a common task in deep neural networks, where the goal is to predict the class of an input based on its features. Here's an example of how classification can be performed in a deep neural network using the popular MNIST dataset of handwritten digits.

The MNIST dataset contains 60,000 training images and 10,000 testing images of handwritten digits from 0 to 9. Each image is a grayscale 28x28 pixel image, and the task is to classify each image into one of the 10 classes corresponding to the 10 digits.

We can use a convolutional neural network (CNN) to classify the MNIST dataset. A CNN is a type of deep neural network that is commonly used for image classification tasks.

What us CNN-

Convolutional Neural Networks (CNNs) are commonly used for image classification tasks, and they are designed to automatically learn and extract features from input images. Let's consider an example of using a CNN to classify images of handwritten digits.

Department of Computer Engineering,

In a typical CNN architecture for image classification, there are several layers, including convolutional layers, pooling layers, and fully connected layers. Here's a diagram of a simple CNN architecture for the digit classification task:

The input to the network is an image of size 28x28 pixels, and the output is a probability distribution over the 10 possible digits (0 to 9).

The convolutional layers in the CNN apply filters to the input image, looking for specific patterns and features. Each filter produces a feature map that highlights areas of the image that match the filter. The filters are learned during training, so the network can automatically learn which features are most relevant for the classification task.

The pooling layers in the CNN downsample the feature maps, reducing the spatial dimensions of the data. This helps to reduce the number of parameters in the network, while also making the features more robust to small variations in the input image.

The fully connected layers in the CNN take the flattened output from the last pooling layer and perform a classification task by outputting a probability distribution over the 10 possible digits.

During training, the network learns the optimal values of the filters and parameters by minimizing a loss function. This is typically done using stochastic gradient descent or a similar optimization algorithm.

Once trained, the network can be used to classify new images by passing them through the network and computing the output probability distribution.

Overall, CNNs are powerful tools for image recognition tasks and have been used successfully in many applications, including object detection, face recognition, and medical image analysis.

CNNs have a wide range of applications in various fields, some of which are:

Image classification: CNNs are commonly used for image classification tasks, such as identifying objects in images and recognizing faces.

Object detection: CNNs can be used for object detection in images and videos, which involves identifying the location of objects in an image and drawing bounding boxes around them.

Semantic segmentation: CNNs can be used for semantic segmentation, which involves partitioning an image into segments and assigning each segment a semantic label (e.g., "road", "sky", "building").

Natural language processing: CNNs can be used for natural language processing tasks, such as sentiment analysis and text classification.

Medical imaging: CNNs are used in medical imaging for tasks such as diagnosing diseases from X-rays and identifying tumors from MRI scans.

Autonomous vehicles: CNNs are used in autonomous vehicles for tasks such as object detection and lane detection.

Department of Computer Engineering,

Video analysis: CNNs can be used for tasks such as video classification, action recognition, and video captioning.

Overall, CNNs are a powerful tool for a wide range of applications, and they have been used successfully in many areas of research and industry.

How Deep Neural Network Work on Classification using CNN-

Deep neural networks using CNNs work on classification tasks by learning to automatically extract features from input images and using those features to make predictions. Here's how it works:

Input layer: The input layer of the network takes in the image data as input.

Convolutional layers: The convolutional layers apply filters to the input images to extract relevant features. Each filter produces a feature map that highlights areas of the image that match the filter.

Activation functions: An activation function is applied to the output of each convolutional layer to introduce non-linearity into the network.

Pooling layers: The pooling layers downsample the feature maps to reduce the spatial dimensions of the data.

Dropout layer: Dropout is used to prevent overfitting by randomly dropping out a percentage of the neurons in the network during training.

Fully connected layers: The fully connected layers take the flattened output from the last pooling layer and perform a classification task by outputting a probability distribution over the possible classes.

Softmax activation function: The softmax activation function is applied to the output of the last fully connected layer to produce a probability distribution over the possible classes.

Loss function: A loss function is used to compute the difference between the predicted probabilities and the actual labels.

Optimization: An optimization algorithm, such as stochastic gradient descent, is used to minimize the loss function by adjusting the values of the network parameters.

Training: The network is trained on a large dataset of labeled images, adjusting the values of the parameters to minimize the loss function.

Prediction: Once trained, the network can be used to classify new images by passing them through the network and computing the output probability distribution.

MNIST Dataset-

The MNIST Fashion dataset is a collection of 70,000 grayscale images of 28x28 pixels, representing 10 different categories of clothing and accessories. The categories include T-shirts/tops, trousers, pullovers, dresses, coats, sandals, shirts, sneakers, bags, and ankle boots.

The dataset is often used as a benchmark for testing image classification algorithms, and it is considered a more challenging version of the original MNIST dataset which contains handwritten digits. The

MNIST Fashion dataset was released by Zalando Research in 2017 and has since become a popular dataset in the machine learning community.

he MNIST Fashion dataset is a collection of 70,000 grayscale images of 28x28 pixels each. These images represent 10 different categories of clothing and accessories, with each category containing 7,000 images.

The categories are as follows:

T-shirt/tops

Trousers

Pullovers

Dresses

Coats

Sandals

Shirts

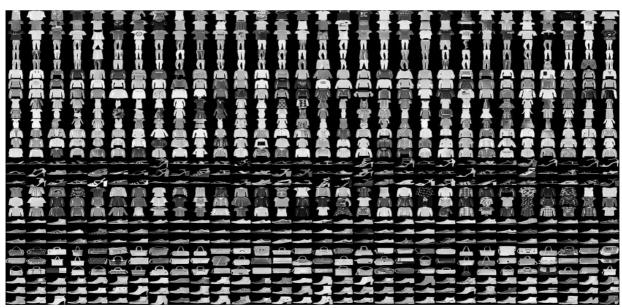
Sneakers

Bags

Ankle boots

The images were obtained from Zalando's online store and are preprocessed to be normalized and centered. The training set contains 60,000 images, while the test set contains 10,000 images. The goal of the dataset is to accurately classify the images into their respective categories.

The MNIST Fashion dataset is often used as a benchmark for testing image classification algorithms, and it is considered a more challenging version of the original MNIST dataset which contains handwritten digits. The dataset is widely used in the machine learning community for research and educational purposes.



Department of Computer Engineering,

Here are the general steps to perform Convolutional Neural Network (CNN) on the MNIST Fashion dataset:

- Import the necessary libraries, including TensorFlow, Keras, NumPy, and Matplotlib.
- Load the dataset using Keras' built-in function, keras.datasets.fashion_mnist.load_data(). This will provide the training and testing sets, which will be used to train and evaluate the CNN.
- Preprocess the data by normalizing the pixel values between 0 and 1, and reshaping the images to be of size (28, 28, 1) for compatibility with the CNN.
- Define the CNN architecture, including the number and size of filters, activation functions, and pooling layers. This can vary based on the specific problem being addressed.
- Compile the model by specifying the loss function, optimizer, and evaluation metrics. Common choices include categorical cross-entropy, Adam optimizer, and accuracy metric.
- Train the CNN on the training set using the fit() function, specifying the number of epochs and batch size.
- Evaluate the performance of the model on the testing set using the evaluate() function. This will provide metrics such as accuracy and loss on the test set.
- Use the trained model to make predictions on new images, if desired, using the predict() function.

Source Code with Output-

```
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow import keras
import numpy as np

(x_train, y_train), (x_test, y_test) = keras.datasets.fashion_mnist.load_data()
```

There are 10 image classes in this dataset and each class has a mapping corresponding to the following labels:

```
#0 T-shirt/top
```

#1 Trouser

#2 pullover

#3 Dress

#4 Coat

#5 sandals

Department of Computer Engineering,

#6 shirt

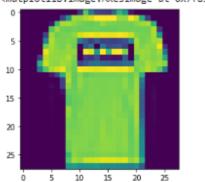
#7 sneaker

#8 bag

#9 ankle boot

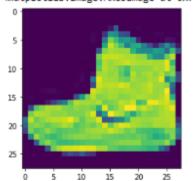
plt.imshow(x_train[1])

<matplotlib.image.AxesImage at 0x7f85874f3a00>



plt.imshow(x_train[0])

<matplotlib.image.AxesImage at 0x7f8584b93d00>



Next, we will preprocess the data by scaling the pixel values to be between 0 and 1, and then reshaping the images to be 28x28 pixels.

 $x_{train} = x_{train.astype('float32') / 255.0}$ $x_{test} = x_{test.astype('float32') / 255.0}$

 $x_{train} = x_{train.reshape}(-1, 28, 28, 1)$ $x_{test} = x_{test.reshape}(-1, 28, 28, 1)$

28, 28 comes from width, height, 1 comes from the number of channels

-1 means that the length in that dimension is inferred.

This is done based on the constraint that the number of elements in an ndarray or Tensor when reshaped must remain the same.

Department of Computer Engineering,

```
# each image is a row vector (784 elements) and there are lots of such rows (let it be n, so there are 784n
   elements). So TensorFlow can infer that -1 is n.
   # converting the training images array to 4 dimensional array with sizes 60000, 28, 28, 1 for 0th to 3rd
   dimension.
   x train.shape
   (60000, 28, 28)
   x test.shape
   (10000, 28, 28, 1)
   y_train.shape
   (60000,)
   y_test.shape
   (10000,)
   # We will use a convolutional neural network (CNN) to classify the fashion items.
   # The CNN will consist of multiple convolutional layers followed by max pooling,
   # dropout, and dense layers. Here is the code for the model:
   model = keras.Sequential([
     keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)),
     # 32 filters (default), randomly initialized
     # 3*3 is Size of Filter
     # 28,28,1 size of Input Image
     # No zero-padding: every output 2 pixels less in every dimension
     # in Paramter shwon 320 is value of weights: (3x3 filter weights + 32 bias) * 32 filters
     # 32*3*3=288(Total)+32(bias)= 320
     keras.layers.MaxPooling2D((2,2)),
     # It shown 13 * 13 size image with 32 channel or filter or depth.
     keras.layers.Dropout(0.25),
     # Reduce Overfitting of Training sample drop out 25% Neuron
     keras.layers.Conv2D(64, (3,3), activation='relu'),
     # Deeper layers use 64 filters
     # 3*3 is Size of Filter
     # Observe how the input image on 28x28x1 is transformed to a 3x3x64 feature map
     # 13(Size)-3(Filter Size)+1(bias)=11 Size for Width and Height with 64 Depth or filter or channel
     # in Paramter shwon 18496 is value of weights: (3x3 filter weights + 64 bias) * 64 filters
     # 64*3*3=576+1=577*32 + 32(bias)=18496
   keras.layers.MaxPooling2D((2,2)),
     # It shown 5 * 5 size image with 64 channel or filter or depth.
   keras.layers.Dropout(0.25),
Department of Computer Engineering,
```

```
keras.layers.Conv2D(128, (3,3), activation='relu'),
     # Deeper layers use 128 filters
     # 3*3 is Size of Filter
     # Observe how the input image on 28x28x1 is transformed to a 3x3x128 feature map
      # It show 5(Size)-3(Filter Size )+1(bias)=3 Size for Width and Height with 64 Depth or filter or
   channel
     # 128*3*3=1152+1=1153*64 + 64(bias)= 73856
     # To classify the images, we still need a Dense and Softmax layer.
     # We need to flatten the 3x3x128 feature map to a vector of size 1152
      keras.layers.Flatten(),
     keras.layers.Dense(128, activation='relu'),
     # 128 Size of Node in Dense Layer
     # 1152*128 = 147584
     keras.layers.Dropout(0.25),
     keras.layers.Dense(10, activation='softmax')
     # 10 Size of Node another Dense Layer
     # 128*10+10 bias= 1290
   1)
   model.summary()
   Model: "sequential"
   Layer (type)
                         Output Shape
                                              Param #
   conv2d (Conv2D)
                             (None, 26, 26, 32)
                                                   320
   max_pooling2d (MaxPooling2D (None, 13, 13, 32)
                                                         0
   )
   dropout (Dropout)
                           (None, 13, 13, 32)
                                                  0
   conv2d_1 (Conv2D)
                              (None, 11, 11, 64)
                                                    18496
   max_pooling2d_1 (MaxPooling (None, 5, 5, 64)
   2D)
   dropout_1 (Dropout)
                                                  0
                             (None, 5, 5, 64)
Department of Computer Engineering,
```

```
conv2d_2 (Conv2D)
                          (None, 3, 3, 128)
                                               73856
flatten (Flatten)
                     (None, 1152)
                                          0
dense (Dense)
                      (None, 128)
                                          147584
dropout_2 (Dropout)
                         (None, 128)
                                             0
dense_1 (Dense)
                       (None, 10)
                                           1290
Total params: 241,546
```

Trainable params: 241,546 Non-trainable params: 0

Compile and Train the Model

After defining the model, we will compile it and train it on the training data.

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

history = model.fit(x_train, y_train, epochs=10, validation_data=(x_test, y_test))

1875 is a number of batches. By default batches contain 32 samles.60000 / 32 = 1875

Finally, we will evaluate the performance of the model on the test data.

test_loss, test_acc = model.evaluate(x_test, y_test)

print('Test accuracy:', test_acc)

Test accuracy: 0.9031000137329102

Conclusion- In this way we can Classify fashion clothing into categories using CNN.

Assignment Question

- 1. What is Binary Classification?
- 2. What is binary Cross Entropy?
- 3. What is Validation Split?
- 4. What is the Epoch Cycle?
- 5. What is Adam Optimizer?